

# E-Commerce Analytics Hub with Predictive Inventory (AWS)

## Problem Statement (Focused & Scoped for 1 Week)

E-commerce businesses struggle with:

1. **Delayed stockout detection** – Traditional analytics take hours/days to process, missing real-time sales trends.
2. **Manual inventory forecasting** – Teams rely on spreadsheets, leading to overstocking or stockouts.
3. **High AWS costs** – Complex pipelines can exceed free-tier limits if not optimized.

**Goal:** Build a **real-time analytics pipeline** that:

- ✓ Processes live sales data
- ✓ Predicts stockouts using ML
- ✓ Runs entirely on AWS Free Tier
- ✓ Delivers insights in <5 min latency

---

## Solution (1-Week MVP Approach)

### 1. Data Ingestion (Kinesis - Free Tier)

- Simulate real-time sales data (Python script → Kinesis)
- Free Tier: 1 shard @ 1MB/s (enough for demo)

### 2. ETL (Glue PySpark - Serverless & Free Tier)

- Deduplicate, clean, and aggregate sales data
- Feature engineering (rolling sales averages, product trends)
- **Cost Saver:** Use Glue Bookmarks to avoid reprocessing

### 3. Predictive Inventory (Redshift ML - Free Tier)

- Train a **simple demand-forecasting model** (ARIMA or Linear Regression)
- Predict stockouts based on sales velocity
- Free Tier: Redshift Serverless (up to 8 hours/day)

### 4. Serverless Alerting (Lambda + SNS - Free Tier)

- Trigger alerts when stockout risk >80%
- Send email/SMS via SNS (Free Tier: 1M requests/month)

### 5. Dashboard (QuickSight - Free Tier)

- Visualize real-time sales vs. predictions
- Highlight at-risk products

# Repository Structure (1-Week MVP)

text

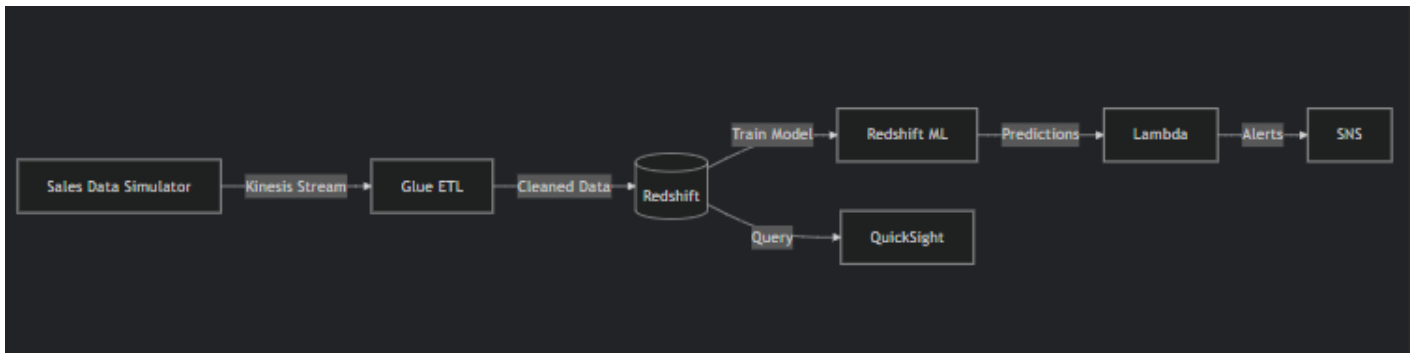
 Copy  Download

```
ecom-stockout-predictor/
|
├── **data-simulator/**
|   └── kinesis-producer.py # Simulate sales data (Python + Boto3)
|
├── **glue-etl/**
|   ├── sales-cleaning.py # Glue PySpark job (dedupe, aggregate)
|   └── feature-engineering.py # Rolling averages, trend features
|
├── **redshift-ml/**
|   ├── train-model.sql # Redshift ML training script
|   └── predict-stockouts.sql # Daily predictions
|
├── **lambda-alerts/**
|   └── stockout-alert.py # Check predictions → SNS alert
|
├── **terraform/** # Infrastructure-as-Code (optional)
|   └── main.tf # AWS setup (Kinesis, Glue, Redshift)
|
├── **quicksight-dash/** # JSON export of dashboard
|   └── stockout-dashboard.json
|
└── **README.md** # Setup guide, Free Tier limits
```

ecom-stockout-predictor/

```
|
|
├── **data-simulator/**
|   └── kinesis-producer.py # Simulate sales data (Python + Boto3)
|
|
├── **glue-etl/**
|   ├── sales-cleaning.py # Glue PySpark job (dedupe, aggregate)
|   └── feature-engineering.py # Rolling averages, trend features
|
|
├── **redshift-ml/**
|   ├── train-model.sql # Redshift ML training script
|   └── predict-stockouts.sql # Daily predictions
|
|
├── **lambda-alerts/**
|   └── stockout-alert.py # Check predictions → SNS alert
```

```
|
|└─ **terraform/**      # Infrastructure-as-Code (optional)
|  └─ main.tf          # AWS setup (Kinesis, Glue, Redshift)
|
|└─ **quicksight-dash/** # JSON export of dashboard
|  └─ stockout-dashboard.json
|
|└─ **README.md**      # Setup guide, Free Tier limits
```









Search

[Alt+S]

United States (N. Virginia)

Getting started

ETL jobs

Visual ETL

Notebooks

Job run monitoring

Data Catalog tables

Data connections

Workflows (orchestration)

Zero-ETL integrations

Data Catalog

Data Integration and ETL

Legacy pages

What's New

Documentation

AWS Marketplace

Enable compact mode

Enable new navigation

Successfully started job

ecom-sales-cleaning

VisualScriptJob detailsRunsData qualitySchedulesVersion Control

Job runs (1/1)

Run statusRetriesStart time (Local)End time (Local)DurationCapacit...Worker typeGlue version

Running008/11/2025 15:58:59-2 m 29 s2 DPUsg.t1xS.0

Run detailsInput arguments (9)LogsRun insightsMetricsTroubleshooting analysis - previewSpark UI

Logs info

```
Driver logs
{
  "metadata": {
    "streamName": "ecom-sales-stream",
    "batchId": "17"
  },
  "shardId-000000000001": {
    "iteratorType": "AFTER_SEQUENCE_NUMBER",
    "iteratorPosition": "496660291558129579571655508148531822025416979908868098"
  },
  "shardId-000000000002": {
    "iteratorType": "AFTER_SEQUENCE_NUMBER",
    "iteratorPosition": "4966602915499065721196702018640459015387329835988854786"
  },
  "shardId-000000000003": {
    "iteratorType": "AFTER_SEQUENCE_NUMBER",
    "iteratorPosition": "49666029155857559447562612846496289981266384948666185986"
  },
  "latestOffset": {
    "shardId-000000000002": {
```

Amazon S3

General purpose buckets

Directory buckets

Table buckets

Vector buckets

Access Grants

Access Points (General Purpose Buckets, FSx file systems)

Access Points (Directory Buckets)

Object Lambda Access Points

Multi-Region Access Points

Batch Operations

IAM Access Analyzer for S3

Block Public Access settings for this account

General purpose buckets

Directory buckets

General purpose buckets (2)

Find buckets by name

Name	AWS Region	Creation date
aws-glue-assets-183900808320-us-east-1	US East (N. Virginia) us-east-1	August 11, 2025, 12:32:23 (UTC+05:30)
ecom-glue-temp	US East (N. Virginia) us-east-1	August 11, 2025, 11:49:47 (UTC+05:30)

Account snapshot

External access summary - new

Amazon S3

General purpose buckets

Directory buckets

Table buckets

Vector buckets

Access Grants

Access Points (General Purpose Buckets, FSx file systems)

Access Points (Directory Buckets)

Object Lambda Access Points

Multi-Region Access Points

Batch Operations

IAM Access Analyzer for S3

Block Public Access settings for this account

Storage Lens

Dashboards

ecom-glue-temp

ObjectsMetadataPropertiesPermissionsMetricsManagementAccess Points

Objects (3)

Find objects by prefix

Name	Type	Last modified	Size	Storage class
cleaned_data/	Folder	-	-	-
scripts/	Folder	-	-	-
temp/	Folder	-	-	-

Amazon S3

General purpose buckets

Directory buckets

Table buckets

Vector buckets [Preview](#)

Access Grants

Access Points (General Purpose Buckets, FSx file systems)

Access Points (Directory Buckets)

Object Lambda Access Points

Multi-Region Access Points

Batch Operations

IAM Access Analyzer for S3

Block Public Access settings for this account

Storage Lens

Dashboards

Storage Lens groups

AWS Organizations settings

Feature spotlight 11

AWS Marketplace for S3

checkpoint/

Objects

Properties

Objects (7)

Copy S3 URI

Copy URL

Download

Open

Delete

Actions

Create folder

Upload

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Find objects by prefix

< 1 >

<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class
<input type="checkbox"/>	commits_\$folder\$	-	August 11, 2025, 12:37:38 (UTC+05:30)	0 B	Standard
<input type="checkbox"/>	commits/	Folder	-	-	-
<input type="checkbox"/>	metadata	-	August 11, 2025, 12:37:38 (UTC+05:30)	45.0 B	Standard
<input type="checkbox"/>	offsets_\$folder\$	-	August 11, 2025, 12:37:38 (UTC+05:30)	0 B	Standard
<input type="checkbox"/>	offsets/	Folder	-	-	-
<input type="checkbox"/>	sources_\$folder\$	-	August 11, 2025, 12:37:45 (UTC+05:30)	0 B	Standard
<input type="checkbox"/>	sources/	Folder	-	-	-

Amazon S3

General purpose buckets

Directory buckets

Table buckets

Vector buckets [Preview](#)

Access Grants

Access Points (General Purpose Buckets, FSx file systems)

Access Points (Directory Buckets)

Object Lambda Access Points

Multi-Region Access Points

Batch Operations

IAM Access Analyzer for S3

Block Public Access settings for this account

Storage Lens

Dashboards

Storage Lens groups

AWS Organizations settings

Feature spotlight 11

AWS Marketplace for S3

ingest\_hour=10/

Objects

Properties

Objects (9)

Copy S3 URI

Copy URL

Download

Open

Delete

Actions

Create folder

Upload

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Find objects by prefix

< 1 >



<input type="checkbox"/>	Name	Type	Last modified	Size	Storage class
<input type="checkbox"/>	part-1754908196744-000000-snappy.parquet	parquet	August 11, 2025, 15:59:59 (UTC+05:30)	2.1 KB	Standard
<input type="checkbox"/>	part-1754908196744-000001-snappy.parquet	parquet	August 11, 2025, 15:59:59 (UTC+05:30)	2.0 KB	Standard
<input type="checkbox"/>	part-1754908196744-000002-snappy.parquet	parquet	August 11, 2025, 15:59:59 (UTC+05:30)	1.8 KB	Standard
<input type="checkbox"/>	part-1754908196744-000003-snappy.parquet	parquet	August 11, 2025, 15:59:59 (UTC+05:30)	1.9 KB	Standard
<input type="checkbox"/>	part-17549082026118-000001-snappy.parquet	parquet	August 11, 2025, 16:00:04 (UTC+05:30)	1.3 KB	Standard
<input type="checkbox"/>	part-17549082026118-000002-snappy.parquet	parquet	August 11, 2025, 16:01:43 (UTC+05:30)	1.3 KB	Standard
<input type="checkbox"/>	part-17549082026118-000003-snappy.parquet	parquet	August 11, 2025, 16:01:43 (UTC+05:30)	1.3 KB	Standard
<input type="checkbox"/>	part-17549082026118-000004-snappy.parquet	parquet	August 11, 2025, 16:01:43 (UTC+05:30)	1.1 KB	Standard



AWS Glue

- Getting started
- ETL jobs
  - Visual ETL
  - Notebooks
  - Job run monitoring
- Data Catalog tables
  - Data connections
  - Workflows (orchestration)
  - Zero-ETL integrations New
- Data Catalog
  - Databases
    - Tables
  - Stream schema registries
  - Schemas
  - Connections
  - Crawlers
  - Classifiers
  - Catalog settings
- Data Integration and ETL
- Legacy pages

What's New   
Documentation   
AWS Marketplace

-  Enable compact mode
-  Enable new navigation

 Announcing new optimization features for Apache Iceberg tables  
Optimize storage for Apache Iceberg tables with automatic snapshot retention and orphan file deletion. [Learn more](#) 

cleaned\_data



Last updated 6/7/23  
August 12, 2025 at 03:57:42  Version 1 (Current version)  [Actions](#)

Table overview

Data quality - new

Table details

Name

cleaned\_data

Database

ecom\_db

Description

-

Last updated

August 11, 2025 at 14:50:45

Classification

Parquet

Location

s3://ecom-glue-temp/cleaned\_data/

Connection

-

Deprecated

-

Column statistics

No statistics

Advanced properties

Schema

Partitions

Indexes

Column statistics - new

Schema (9)

View and manage the table schema.

[Edit schema as JSON](#)

[Edit schema](#)

 Filter schemas

#	Column name	Data type	Partition key	Comment
1	order_id	string	-	-
2	product_id	int	-	-
3	quantity	int	-	-
4	price	double	-	-
5	timestamp	string	-	-
6	ingest_year	string	Partition (0)	-
7	ingest_month	string	Partition (1)	-
8	ingest_day	string	Partition (2)	-
9	ingest_hour	string	Partition (3)	-

- ☒ Create a new namespace
- ☐ Add to an existing namespace

#### Namespace

This is a unique name that defines the namespace.

The name must be from 3-64 characters. Valid characters are a-z (lowercase only), 0-9 (numbers), and - (hyphen).

### Database name and password

#### Database name

The name of the first database in the Amazon Redshift Serverless environment.

The name must be 1-64 alphanumeric characters (lowercase only), and it can't be a reserved word.

#### Admin user credentials

IAM credentials provided as your default admin user credentials. To add a new admin username and password, customize admin user credentials.

- ☒ **Customize admin user credentials**  
To use the default IAM credentials, clear this option.

#### Admin user name

The administrator's user name for the first database.

The name must be 1-128 alphanumeric characters, and it can't be a reserved word.

#### Admin password

Select an option to manage your admin password.

- ☐ **Manage admin credentials in AWS Secrets Manager** [Info](#)  
AWS manages a KMS key that encrypts your data.
- ☐ **Generate a password**  
Amazon Redshift generates an admin password.
- ☒ **Manually add the admin password**  
Manually enter the admin password.


#### Admin user password

The password of the admin user.

- ☒ **Show password**

Must be 8-64 characters long. Must contain at least one uppercase letter, one lowercase letter and one number. Can be any printable ASCII character except "/", "", or "@".

### Permissions

-  Associate an IAM role so that your serverless endpoint can LOAD and UNLOAD data. You can create an IAM role or the default for this configuration that has the `AmazonRedshiftAllCommandsFullAccess` policy.

Redshift\_access\_S3

Info

Delete

Allows Redshift clusters to call AWS services on your behalf.

Summary

Edit

Creation date

August 11, 2025, 17:05 (UTC+05:30)

Last activity

41 minutes ago

ARN

arn:aws:iam::183900808320:role/Redshift\_access\_S3

Maximum session duration

1 hour

Permissions

Trust relationships

Tags

Last Accessed

Revoke sessions

Permissions policies (8)

Info

Simulate

Remove

Add permissions

You can attach up to 10 managed policies.

Q Search

Filter by Type

All types

< 1 >

<input type="checkbox"/>	Policy name	Type	Attached entities
<input type="checkbox"/>	<div><div></div>AmazonKinesisFullAccess</div>	AWS managed	3
<input type="checkbox"/>	<div><div></div>AmazonRedshiftDataFullAccess</div>	AWS managed	1
<input type="checkbox"/>	<div><div></div>AmazonS3ReadOnlyAccess</div>	AWS managed	1
<input type="checkbox"/>	<div><div></div>AWSGlueConsoleFullAccess</div>	AWS managed	1
<input type="checkbox"/>	<div><div></div>AwsGlueDataBrewFullAccessPolicy</div>	AWS managed	1
<input type="checkbox"/>	<div><div></div>AWSGlueSchemaRegistryFullAccess</div>	AWS managed	1
<input type="checkbox"/>	<div><div></div>AWSLambda_FullAccess</div>	AWS managed	1
<input type="checkbox"/>	<div><div></div>S3_glue</div>	Customer inline	0

► Permissions boundary (not set)

▼ Generate policy based on CloudTrail events

You can generate a new policy based on the access activity for this role, then customize, create, and attach it to this role. AWS uses your CloudTrail events to identify the services and actions used and generate a policy. [Learn more](#)

Generate policy

No requests to generate a policy in the past 7 days.

IAM > Roles > Redshift\_access\_S3

Info

Identity and Access Management (IAM)

<

Q Search IAM

Dashboard

▼ Access management

User groups

Users

Roles

Policies

Identity providers

Account settings

Root access management 

New

▼ Access reports

Access Analyzer

Resource analysis 

New

Unused access

Analyzer settings

Credential report

Organization activity

Service control policies

Resource control policies 

New

IAM Identity Center

AWS Organizations

Policy was successfully attached to role.

X

Redshift\_access\_S3

Info

Delete

Allows Redshift clusters to call AWS services on your behalf.

Summary

Edit

Creation date

August 11, 2025, 17:05 (UTC+05:30)

Last activity

-

ARN

arn:aws:iam::183900808320:role/Redshift\_access\_S3

Maximum session duration

1 hour

Permissions

Trust relationships

Tags

Last Accessed

Revoke sessions

Permissions policies (4)

Info

Simulate

Remove

Add permissions

You can attach up to 10 managed policies.

Q Search

Filter by Type

All types

< 1 >

<input type="checkbox"/>	Policy name	Type	Attached entities
<input type="checkbox"/>	<div><div></div>AmazonKinesisFullAccess</div>	AWS managed	3
<input type="checkbox"/>	<div><div></div>AmazonS3ReadOnlyAccess</div>	AWS managed	1
<input type="checkbox"/>	<div><div></div>AWSGlueConsoleFullAccess</div>	AWS managed	1
<input type="checkbox"/>	<div><div></div>AWSLambda_FullAccess</div>	AWS managed	1

aws

Search

[Alt+S]

United States (N. Virginia)

samr

✓

The namespace ecom-namespace on Amazon Redshift has been successfully registered and the invitation is pending approval. In order to be connected to the AWS Glue Data Catalog a Redshift federated catalog will need to be created after the invitation is approved.

Go to AWS Glue Data Catalog

✕

Amazon Redshift Serverless > Namespace configuration > ecom-namespace

ecom-namespace

Info

Actions

Query data

General information

Namespace ecom-namespace	Status ✓ Available	Admin user name admin
Namespace ID d26ba560-8e51-4efc-a4e9-d7ad6aade4b8	Date created August 11, 2025, 17:08 (UTC+05:30)	Database name dev
Namespace ARN arn:aws:redshift-serverless:us-east-1:183900808320:namespace/d26ba560-8e51-4efc-a4e9-d7ad6aade4b8	Total used storage 75 MB	Total table count -
Namespace register status ✓ Registered to AWS Glue Data Catalog		

Workgroup

Data backup

Database

Security and encryption

Datashares

Zero-ETL integrations

Resource policy

Tags

Workgroup name

Actions

Set up compute resources for your workgroup.

Workgroup ecom-analytics	Status ✓ Available
-----------------------------	-----------------------

Used a crawler to read the schema

AWS Glue

Crawlers

S3-reader

AWS Glue

Getting started

ETL jobs

Visual ETL

Notebooks

Job run monitoring

Data Catalog tables

Data connections

Workflows (orchestration)

Zero-ETL integrations

New

Data Catalog

Databases

Tables

Stream schema registries

Schemas

Connections

Crawlers

Classifiers

Catalog settings

Data Integration and ETL

Legacy pages

What's New

Documentation

AWS Marketplace

Enable compact mode

Enable new navigation

Crawler successfully starting

The following crawler is now starting: "S3-reader"

Last updated (UTC)

August 11, 2025 at 14:50:57

Run crawler

Edit

Delete

S3-reader

Crawler properties

Name	S3-reader	IAM role	BlueKinesisRole	Database	ecom_db	State	STOPPING
Description	-	Security configuration	-	Lake Formation configuration	-	Table prefix	-
Maximum table threshold	-						

Advanced settings

Crawler runs

Schedule

Data sources

Classifiers

Tags

Crawler runs (1)

Stop run

View CloudWatch logs

View run details

The list of crawler runs for this crawler.

Filter data

Filter by a date and time range

Start time (UTC)

End time (UTC)

Current/last duration

Status

DPU hours

August 11, 2025 at 14:4...

August 11, 2025 at 14:5...

01 min 06 s

Completed

-

Saved it to ecom\_db database and deleted crawler

Created the table

aws

Search

[Alt+S]

United States (N. Virginia)

Account ID: 1839-0080-8320

samrai-23

Editor

Queries

Notebooks

Charts

History

Scheduled queries

Redshift query editor v2

Create

Load data

Filter resources

Serverless: ecom-analytics-wg

native databases (2)

dev

ecom

Tables

1

sales\_features

0

Views

0

Functions

0

Stored procedures

0

public

sample\_data\_dev

external databases (1)

Untitled 1

Run

Limit 100

Explain

Isolated session

Serverless: ec...

dev

Schedule

1

2

3

4

5

6

7

8

9

10

11

12

13

CREATE TABLE ecom.sales\_features (

order\_id VARCHAR(255),

product\_id INTEGER,

quantity INTEGER,

price DOUBLE PRECISION,

timestamp VARCHAR(255), -- Keep as string if format isn't standard timestamp

ingest\_year VARCHAR(4),

ingest\_month VARCHAR(2),

ingest\_day VARCHAR(2),

ingest\_hour VARCHAR(2)

)

DISTKEY(product\_id)

SORTKEY(timestamp, ingest\_year, ingest\_month);

Row 13, Col 47, Chr 409

Result 1

Export

Chart

Summary

Returned rows: 0

Elapsed time: 317ms

Result set query:

/\* RQEV2-GcrHC7LS4C \*/

CREATE TABLE ecom.sales\_features (

order\_id VARCHAR(255),

product\_id INTEGER,

quantity INTEGER,

price DOUBLE PRECISION,

timestamp VARCHAR(255), -- Keep as string if format isn't standard timestamp

ingest\_year VARCHAR(4),

ingest\_month VARCHAR(2),

ingest\_day VARCHAR(2),

ingest\_hour VARCHAR(2)

)

DISTKEY(product\_id)

SORTKEY(timestamp, ingest\_year, ingest\_month)

Redshift query editor v2

Filter resources

Serverless: ecom-analytics-wg

- native databases (2)
  - dev
    - ecom
    - public
    - spectrum
      - Tables
        - cleaned\_data 1
    - temp\_spectrum
    - sample\_data\_dev
  - external databases (1)

```
1 -- Check row count
2 SELECT COUNT(*) FROM spectrum.cleaned_data;
3
4 -- Sample 10 records
5 SELECT * FROM spectrum.cleaned_data LIMIT 10;
6
7 -- Check for NULLs in key columns
8 SELECT
9   COUNT(*) FILTER (WHERE product_id IS NULL) AS null_products,
10  COUNT(*) FILTER (WHERE quantity IS NULL) AS null_quantities
11 FROM spectrum.cleaned_data;
```

Run Limit 100 Explain Isolated session Serverless: ec... dev

Result 1 (1)

count
87

Redshift query editor v2

Filter resources

Serverless: ecom-analytics-wg

- native databases (2)
  - dev
    - ecom
    - public
    - spectrum
      - Tables
        - cleaned\_data 1
    - temp\_spectrum
    - sample\_data\_dev
  - external databases (1)

```
1 -- Check row count
2 SELECT COUNT(*) FROM spectrum.cleaned_data;
3
4 -- Sample 10 records
5 SELECT * FROM spectrum.cleaned_data LIMIT 10;
6
7 -- Check for NULLs in key columns
8 SELECT
9   COUNT(*) FILTER (WHERE product_id IS NULL) AS null_products,
10  COUNT(*) FILTER (WHERE quantity IS NULL) AS null_quantities
11 FROM spectrum.cleaned_data;
```

Run Limit 100 Explain Isolated session Serverless: ec... dev

Result 1 (1) Result 2 (10) Result 3

order id	product id	quantity	price	timestamp	ingest year	ingest month	ingest day	ingest hour
01e00e64-7168-42e6-ac0...	94	3	31.97	2025-06-01T21:51:46	2025	06	11	10
0e7cd3af-aaa6-495f-bd8e...	33	1	35.18	2025-06-05T13:31:26	2025	06	11	10
11679875-4739-4206-891...	40	3	28.01	2025-06-01T07:47:36	2025	06	11	10
18b5dbf1-b169-499a-850...	25	3	82.6	2025-06-10T19:38:06	2025	06	11	10
44c1bbdb-94df-44b3-841...	37	2	25.23	2025-06-08T05:18:48	2025	06	11	10
4804a869-67eb-4411-87c...	62	1	78.54	2025-06-04T10:02:54	2025	06	11	10
4a961e87-8d77-4f64-a26...	50	6	65.05	2025-06-05T07:24:54	2025	06	11	10
616209b0-b6ce-4b74-b2b...	47	5	77.06	2025-06-05T08:22:32	2025	06	11	10
7f2eac3a-2a65-4a85-b357...	44	1	11.64	2025-06-02T21:55:39	2025	06	11	10
7af29b05-c2ce-4736-874...	47	2	72.15	2025-06-05T05:54:54	2025	06	11	10

Redshift query editor v2

Create Load data

Filter resources

Serverless: ecom-analytics-wg

- native databases (2)
  - dev
    - ecom
    - public
    - spectrum
      - Tables 1
        - cleaned\_data
      - temp\_spectrum
      - sample\_data\_dev
  - external databases (1)

Run Limit 100 Explain Isolated session Serverless: ec... dev

```
1
2 -- Check for NULLs in key columns
3 SELECT
4   SUM(CASE WHEN product_id IS NULL THEN 1 ELSE 0 END) AS null_products,
5   SUM(CASE WHEN quantity IS NULL THEN 1 ELSE 0 END) AS null_quantities
6 FROM spectrum.cleaned_data;
```

Result 1 (1)

null products	null quantities
0	0

CloudWatch Alarms Bill alarm

Alarms (1)

Search

Alarm state: Any

Alarm type: Any

Actions status: Any

Hide Auto Scaling alarms

Bill alarm

Metric alarm

In alarm

Graph

EstimatedCharges

EstimatedCharges > 1 for 1 datapoints within 6 hours

No unit

4.4

2.2

0

08/05 08/06 08/07 08/08 08/09 08/10 08/11 08/12

EstimatedCharges

Click timeline to see the state change at the selected time.

In alarm OK Insufficient data Disabled actions

Details Tags Actions History Parent alarms

Details

Name

Bill alarm

Type

Metric alarm

Description

No description

State

In alarm

Threshold

EstimatedCharges > 1 for 1 datapoints within 6 hours

Last state update

2025-08-12 05:53:56 (Local)

Actions

Actions enabled

Namespace

AWS/Billing

Metric name

EstimatedCharges

Currency

USD

Statistic

Maximum

Period

6 hours

Datapoints to alarm

1 out of 1

Missing data treatment

Treat missing data as missing

Percentiles with low samples

evaluate

ARN

arn:aws:cloudwatch:us-east-1:183900808320:alarm:Bill alarm



Cost summary info

Month-to-date cost  
**\$4.39**  
- compared to last month for same period

Total forecasted cost for current month  
**Data unavailable**

Last month's cost for same time period  
**\$0.00**  
Jul 1 – 12

Last month's total cost  
**\$0.00**

[View bill](#)


Cost monitor info

Budgets status  
**1 over budget**  
1 active budget(s)  
Cost anomalies status (MTD)  
**None detected**  
1 monitor(s) active

Cost breakdown info

Group costs by  
Service

Costs (\$)



[Analyze your costs in Cost Explorer](#)

Recommended actions (3) info


- Budget status**  
My Zero-Spend Budget is over budget. [Review budgets](#)
- Budget threshold exceeded**  
1 of your budget alert thresholds have been exceeded. [Review budgets](#)
- Getting started**  
Add an additional billing contact. [Update billing contact.](#)

Top trends (4) info

Group by  
Service  
Time range  
Jun 2025 to Jul 2025

Cost allocation coverage info

- Billing and Cost Management
- Choose billing view New  
Primary view
- Home
- Getting Started
- Billing and Payments
- Bills
- Payments
- Credits
- Purchase Orders
- Cost and Usage Analysis
- Cost Explorer
- Cost Explorer Saved Reports
- Cost Anomaly Detection
- Free Tier
- Data Exports
- Customer Carbon Footprint Tool
- Cost Organization
- Cost Categories
- Cost Allocation Tags
- Billing Conductor
- Budgets and Planning
- Budgets New
- Budgets Reports
- Pricing Calculator New
- Savings and Commitments
- Cost Optimization Hub New
- Savings Plans
- Overview
- Inventory



Cost and usage breakdown (11)

Find cost and usage data

[Download as CSV](#)

Service	Service total	August 2025*
Total costs	\$4.39	\$4.39
Glue	\$3.43	\$3.43
Tax	\$0.68	\$0.68
Kinesis	\$0.28	\$0.28
CloudShell	\$0.00	\$0.00
Key Management Service	\$0.00	\$0.00
EC2-Other	\$0.00	\$0.00
EC2-Instances	\$0.00	\$0.00
S3	\$0.00	\$0.00
VPC	\$0.00	\$0.00
CloudWatch	\$0.00	\$0.00

\*Charges for your current billing period shown on these reports are estimated charges. Estimated charges shown on this page, or shown on any notifications that we send to you, may differ from your actual charges for this statement period. This is because estimated charges presented on this page do not include usage charges accrued during this statement.

- VPC dashboard
- EC2 Global View
- Filter by VPC
- Virtual private cloud
- Your VPCs
- Subnets
- Route tables
- Internet gateways
- Egress-only internet gateways

Your VPCs (1) info

Find VPCs by attribute or tag

Name	VPC ID	State	Block Public...	IPv4 CIDR	IPv6 CIDR	DHCP option set	Main route table	Main network ACL	Tenancy	Default VPC	On...
-	vpc-0663b609c3e7d1ab	Available	Off	172.31.0.0/16	-	dhcp-0d2d0a5f9a2455...	vpc-00a70e230957db7a0	acl-07552ff181e18554	default	Yes	On...